

Distributed Inference on xv6 — Milestone 1: Understand and Characterise

Team samosaOS

September 23, 2026

Contents

Team Information	3
Warm-Up (Section 3.0)	4
Method	4
Call Graph	4
The Three Required Facts, and How Each Was Established	4
Runbook (Section 3.2)	7
Prerequisites	7
1. Clone and install the toolchain	7
2. Fetch the model checkpoints (~495 MB, ~2-3 min)	7
3. Single node	7
4. Weight-fetch path	8
5. Multi-node ring	8
6. Regression suites	8
7. Reverse-engineering the kernel (warm-up + kernel trace)	9
Bring-Up Evidence (Section 3.1)	10
Single Node	10
Multi-Node Ring (2 workers)	10
Weight-Fetch Path	10
Regression Suites	13
Architecture Map (Section 3.3) — Eight Subsystems	15
1. Network Stack (kernel, from binary)	15
2. Remote Procedure Call	15
3. Data Serialisation	15
4. Discovery and Registration	15
5. Ring Driver and Shard Logic	15
6. Transformer Kernels	15
7. Weight Fetch Client	16
8. Kernel Hardening Set (kernel, from binary)	16
Figures (Section 3.4)	18
F1 — System Architecture (flowchart)	18
F2 — Wire Format: RPC Envelope + XDR Activation Message (packet-beta)	19

F3 — Worker Lifecycle (stateDiagram-v2)	22
Where diagrams/state_diagram.png is wrong	22
F4 — One Token, Master → Worker1 → Worker2 → Master (sequenceDiagram)	24
F5 — Inference Hot Path Through user/llama_core.c (flowchart)	25
F6 — Node Memory Map (block-beta)	26
F7 — Baseline Graphs (xychart-beta)	29
Single-node inference time per prompt (ms)	29
Ring generation time: single-node mean vs. 2-worker vs. 3-worker (ms)	29
Per-token latency breakdown, ring only: compute vs. network (µs/token)	31
Difference From Upstream xv6 (Section 3.5)	32
1. Added System Calls	32
2. File Classification	34
Baseline Measurements (Section 3.6)	35
Method	35
Single-Node (XV6_CPUS=1)	35
Multi-Node Ring (first prompt, “Once upon a time”)	36
Observations Log	37
1. server/distinf_sim.py ring runs are flaky under WSL2 even with correct config .	37
2. worker_entry_t.last_seen_ms (user/discovery.h:116) holds ticks, not ms	40
3. discovery_expire_workers ages a PENDING worker exactly like an ACTIVE one — the shipped diagrams/state_diagram.png doesn’t show this and is wrong	40
4. handle_auth_hello’s “slot stays PENDING and ages out” comments (user/discovery.c:514,562) rely on behaviour that isn’t obviously intended	40
5. rdcycle/rdinstret do not measure guest execution under this QEMU configuration — they track host wall-clock time almost exactly	41
Related Systems (Section 3.8)	42
Stock xv6-riscv	42
Linux	42
xv6-llm-runtime-arch (github.com/syedtaha22/xv6-llm-runtime-arch — the base this project builds on)	42
Exo (github.com/exo-explore/exo) — closest prior-art pipeline-parallel inference system	43
Other pipeline-parallel inference systems (general class: Petals, vLLM’s pipeline- parallel mode)	43
Contributions (Section 3.9, bonus)	44
Appendices	45
AI Use Disclosure	45
Repository Layout of Supporting Evidence	45
Commit Log	45

Team Information

Team name: samosaOS

Repository: <https://github.com/Arjcodes123/OS-Milestone-1>

Name	ERP
Abdul Rehman Javaid	30532
Hamza Uzair	30544
Javeria Khan	30565
Qurat Ul Ain	30493

Division of work: all four team members performed the environment setup (cloning the repository, installing the toolchain, and confirming the three configurations boot) independently on their own laptops.

Warm-Up (Section 3.0)

Method

All facts below come from **static inspection only** — the binary was never executed:

```
riscv64-unknown-elf-nm build/challenge | grep -E ' [tT] '  
riscv64-unknown-elf-objdump -d --disassemble=<fn> build/challenge
```

nm gave the complete function list and their addresses:

Address	Function
0x100e8	fib
0x1013e	helper
0x10160	foo
0x101a6	bar
0x101ea	compute

Disassembling each function and reading every jal (jump-and-link) instruction — RISC-V's call instruction — gives every call site and its target:

Caller	Address	Instruction	Target
compute	0x101fe	jal 10160	foo
compute	0x1021a	jal 101a6	bar
bar	0x101c4	jal 10160	foo
bar	0x101d2	jal 10160	foo
foo	0x10174	jal 1013e	helper
foo	0x1018e	jal 100e8	fib
fib	0x10116	jal 100e8	fib
fib	0x10128	jal 100e8	fib
helper	—	(no jal in its body)	leaf, calls nothing

Call Graph

flowchart TD

```
compute --> foo  
compute --> bar  
bar --> foo  
foo --> helper  
foo --> fib  
fib -->|recursive| fib
```

The Three Required Facts, and How Each Was Established

1. **Every function compute reaches.** compute's own disassembly contains exactly two jal targets, foo (0x10160) and bar (0x101a6). Disassembling those in turn shows bar calls foo again (twice), and foo calls helper (0x1013e) and fib (0x100e8). No other jal targets appear anywhere in this call tree, so the complete reachable set from compute is {foo, bar, helper, fib}.
2. **The recursive function.** fib's disassembly contains two jal 100e8 instructions (0x10116 and 0x10128) — 100e8 is fib's own entry address from the nm table, i.e. fib




Figure 1: Recovered call graph for `compute()`

calls itself, twice, from within its own body. That is a direct, statically-visible self-call, which is the definition of recursion at the binary level (consistent with a naive `fib(n-1) + fib(n-2)` implementation, though the source is not available to confirm the body beyond the two self-calls).

3. **The function called from two different places.** Cross-referencing every `jal` target against its caller shows `foo` is the target of calls from **two distinct calling functions**: `compute` (`0x101fe`) and `bar` (`0x101c4` and `0x101d2`). No other non-recursive function is called from more than one caller — `bar` and `fib` are only ever called from `compute` and `foo` respectively, and `helper` is only ever called from `foo`. `foo` is therefore the one function reached from two different places in the graph, distinct from `fib`'s self-recursion.

```

gdb-multiarch: break sys_exec, backtrace, registers, stepi

0x0000000000001000 in ?? ()
Breakpoint 1 at 0x80005856

Breakpoint 1, 0x0000000000005856 in sys_exec ()
#0  0x0000000000005856 in sys_exec ()
#1  0x00000000000014dc in syscall ()
#2  0x0000000000001224 in usertrap ()
#3  0x00000003ffffff124 in ?? ()
a0      0x1      1
a1      0x3ffffbdf8    274877890040
a2      0x1      1
ra      0x800014dc    0x800014dc <syscall+60>
sp      0x3ffffbdf0    0x3ffffbdf0
=> 0x80005856 <sys_exec+30>:  jal    0x8000142c <argaddr>
    0x8000585a <sys_exec+34>:  li      a2,128
    0x8000585e <sys_exec+38>:  addi   a1,s0,-184
    0x80005862 <sys_exec+42>:  li      a0,0
0x000000000000142c in argaddr ()
0x000000000000142e in argaddr ()
0x0000000000001430 in argaddr ()
a0      0x1      1
[Inferior 1 (process 1) detached]

```

Figure 2: Kernel entry-point trace: gdb break at `sys_exec`, backtrace, registers, stepi (Section 2.10 / reverse-engineering workflow)

Raw disassembly evidence backing the call graph above is in `docs/figures/warmup_disassembly_evidence.txt` and the full gdb session in `docs/transcripts/gdb_kernel_trace.txt`.

Runbook (Section 3.2)

What a new team member needs to reproduce all three configurations from a clean checkout. Followed literally, in order, on a clean machine.

Prerequisites

- A Debian/Ubuntu Linux environment. **On Windows, this means WSL2** — native Windows cannot run this toolchain. If using WSL2, one extra step is required beyond a default install (see the callout below).
- git, python3 with venv, and internet access to fetch model checkpoints and apt packages.

WSL2-only prerequisite (skip on native Linux): the multi-node ring configuration uses QEMU multicast networking (-netdev socket,mcast=...) between the host and guest VMs. **WSL2's default NAT networking mode does not forward this traffic** — the ring will hang and fail with never assigned layers within 900s / arp_lookup: resolution timed out. Fix, once, before attempting the ring:

```
# In a Windows PowerShell/cmd prompt (not inside WSL):
notepad %USERPROFILE%\wslconfig
```

Add:

```
[wsl2]
networkingMode=mirrored
```

Then `wsl --shutdown` and reopen a WSL terminal. This is required only for the ring (`distinf_sim.py`); the single-node and weight-fetch configurations work under either networking mode. See Observations Log #1 for the full diagnosis.

1. Clone and install the toolchain

```
git clone <your fork/clone URL>
cd <cloned directory>

sudo apt update
sudo apt install -y git build-essential gdb-multiarch qemu-system-misc \
    gcc-riscv64-linux-gnu binutils-riscv64-unknown-elf python3-venv

python3 -m venv .venv
source .venv/bin/activate
pip install scapy pytest coloredlogs
```

2. Fetch the model checkpoints (~495 MB, ~2-3 min)

```
chmod +x fetch-models.sh boot.sh boot-gdb.sh
./fetch-models.sh
```

Expected: `./fetch-models.sh --check` reports all three files present with matching sizes (`stories15M.bin` 60,816,028 B, `tokenizer.bin` 433,869 B, `stories110M.bin` 438,381,596 B).

3. Single node

```
# terminal 1
source .venv/bin/activate
python server/server.py
```

```
# expect three "Loaded file" lines and "Server listening on 0.0.0.0:9999 (UDP)"

# terminal 2
./boot.sh
# at the xv6 "$ " prompt:
llama -n 32 -i "Once upon a time" -t 0
# exit QEMU: Ctrl-A, release, then x
```

Expected: the first run sits at “Fetching from server...” for 1-2 minutes (first weight fetch over emulated UDP), then prints generated text and a profiling report ending in a SYSTEM PERFORMANCE METRICS block.

4. Weight-fetch path

```
# with server/server.py still running from step 3:
./boot.sh
# at the xv6 prompt:
testftp
```

Expected: two PASS lines — Successfully fetched tokenizer: 433869 bytes and Successfully fetched weights: 60816028 bytes, each with a SHA-256 line.

5. Multi-node ring

Requires the WSL2 mirrored-networking prerequisite above if applicable. Stop server/server.py first — the ring simulator starts its own weight server on the same port.

```
python server/distinf_sim.py --workers 2 --model 1 --steps 32 \
    --prompt "Once upon a time"
```

Expected: [sim] OK lines for boot, registration, layer assignment, and shard residency, then [sim] ==== generated ==== with the same text the single-node run produced, then [sim] PASS. Kill any leftover qemu-system-riscv64 processes before starting a second run — the nodes share one network segment and two runs cannot overlap.

Known flakiness (WSL2 specifically): even with mirrored networking enabled, the ring’s large UDP transfers occasionally lose a small fraction of chunks and the run fails with Transfer failed: received only N/72001 chunks after 6 rounds. This recurred intermittently across repeated identical invocations on this host and is not fixed by retrying the runbook steps differently — see Observations Log #1. A plain retry of the same command usually succeeds.

6. Regression suites

```
pytest server/kernel_regression_test.py -v --noconftest # expect: 14 passed
pytest server/weightfetch_regression_test.py -v --noconftest # expect: 4 passed
pytest server/distinf_regression_test.py -v --noconftest # expect: 2 passed,
↳ 6 skipped
```

On this host, distinf_regression_test.py’s two ring-based tests failed reproducibly (3/3 attempts) with the same environment-level flakiness as step 5, despite a manual distinf_sim.py invocation succeeding — see Observations Log #1.

7. Reverse-engineering the kernel (warm-up + kernel trace)

```
cd xv6-riscv
riscv64-unknown-elf-nm kernel/kernel | grep -E ' [tT] '
riscv64-unknown-elf-objdump -d kernel/kernel
cd ..

# kernel debugger, in one combined shell session:
(./boot-gdb.sh < /dev/null > console.txt 2>&1 &)
sleep 4
gdb-multiarch -batch -ex "target remote localhost:26000" \
    -ex "break sys_exec" -ex "continue" -ex "backtrace" \
    -ex "info registers a0 a1 a2" xv6-riscv/kernel/kernel
```

WSL2 note: run the `boot-gdb.sh` background-launch and the `gdb-multiarch` connect in the *same* shell invocation/session. Launching them in separate terminal sessions that don't share a process group caused the backgrounded QEMU process to receive `SIGHUP` and die before `gdb` could connect, on this host.

Bring-Up Evidence (Section 3.1)

One console transcript per configuration, as booted following the runbook above.

Single Node

Multi-Node Ring (2 workers)

Weight-Fetch Path

```
single-node bring-up: llama -n 32 -i "Once upon a time" -t 0

xv6 kernel is booting

Scheduler: Round Robin (Default)
hart 2 starting
hart 1 starting
llama -n 32 -i "Once upon a time" -t 0
init: starting sh
$ No cached data found for llm_weights. Fetching from server...
PASS    llm_weights fetched and cached in shared memory (ID: 0).

No cached data found for llm_tokenizer. Fetching from server...
PASS    llm_tokenizer fetched and cached in shared memory (ID: 1).

thread_create: sharing user memory with main process (pid=3)
thread_create: sharing user memory with main process (pid=3)
thread_create: sharing user memory with main process (pid=3)
Once upon a time, there was a little girl named Lily. She loved to play outside in the sunshine. One day, she
saw a big

achieved tok/s: 4.525547

=====
FUNCTION PROFILING
=====
-----
Function      Calls Total(ms) Self(ms)  Avg(ms)  Time(%)
-----
fetch_model_weights      1  108040  108040  108040.00  92.79%
  matmul      1376    6154    6267     4.47    5.38%
fetch_tokenizer          1     707     707     707.00    0.60%
  forward        32    6985     523    218.28    0.44%
  encode          1     391     391     391.00    0.33%
  generate        1    7548     138   7548.00    0.11%
multihead_attention     192      90     109      0.46    0.09%
  build_transformer       1      93      93     93.00    0.07%
    rmsnorm      416      47      89      0.11    0.07%
  build_tokenizer         1      35      35     35.00    0.03%
    sample       28      19      23      0.67    0.01%
    decode       32       4       8      0.12    0.00%
  build_sampler          1       3       3      3.00    0.00%
=====
Total Time: 117955 ms

=====
SYSTEM PERFORMANCE METRICS
=====

MEMORY USAGE:
  Initial      1.7 MB (1839104 bytes)
  Peak        66.5 MB (69730304 bytes)
  Final       66.5 MB (69730304 bytes)

TIMING:
  Total (E2EL) 117955 ms
  Inference   6850 ms

LLM METRICS:
  TTFT      112350 ms

THROUGHPUT:
  Generated    31 tokens
  Rate        4.52 tokens/sec
=====

$ QEMU: Terminated
```

Figure 3: Single-node bring-up: generated text and profiling report

```
multi-node ring: distinf_sim.py --workers 2 --model 1

[sim] staging 3 prebuilt node image(s)...
[sim] starting host mcast switch (weight server on the segment)...
[sim] booting 3 node(s) (timeout 240s each)...
[sim] master: expecting 2 worker(s), model 1 (stories15M), 32 tokens
[sim] OK worker1 registered (0.7s)
[sim] OK worker2 registered (1.0s)
[sim] OK 2 worker(s) registered (serialized)
[sim] OK layers assigned
[sim] OK worker2 shard resident (52s)
[sim] OK worker1 shard resident (52s)
[sim] OK all shards resident -- generating

[sim] ==== generated ====
Once upon a time , there was a little girl named Lily . She loved to play outside in the sunshine . One day
, she saw a big
[sim] =====
[sim] ===== run summary =====
[sim] topology      master 10.0.0.2 + 2 worker(s) | model 1 (stories15M) | 32 steps
[sim] build         3 image(s) in 0s
[sim] boot          master 16s | worker1 16s | worker2 17s
[sim] registration  worker1 0.7s | worker2 1.0s   (serialized)
[sim] layers        w1 [0,3) | w2 [3,6)
[sim] shard fetch   worker2 52s | worker1 52s
[sim] final ring    2/2 worker(s)
[sim] generation    32 token(s) in 79.2s
[sim] latency       latency/token (avg over 32): ring rtt 2105377 us = compute 466767 us + network 1638610 us
| master 264496 us | 0.4 tok/s
[sim] rpc drops     malformed 0 | rate-limited 0 | bad-prog 0 | unknown-proc 0
[sim] logs          /tmp/distinf_sim/{master,worker1..N}.log
[sim] result        PASS
[sim] =====
[sim] PASS
```

Figure 4: Ring bring-up: worker registration through generated text

```
weight-fetch path: testftp

xv6 kernel is booting

Scheduler: Round Robin (Default)
hart 1 starting
hart 2 starting
init: starting sh
$ testftp
Starting LLM file transfer test...
PASS Successfully fetched tokenizer: 433869 bytes
Final SHA-256: 50a52ef822ee9e83de5ce9d0be0a025a773d019437f58b5ff9dcafb063ece361
---
PASS Successfully fetched weights: 60816028 bytes
Final SHA-256: cd590644d963867a2b6e5a1107f51fad663c41d79c149fbecbbb1f95fa81f49a
PASS LLM file transfer test completed
$ qemu-system-riscv64: terminating on signal 15 from pid 5561 ()
```

Figure 5: testftp: byte counts for tokenizer and weights

Regression Suites

```
pytest server/kernel_regression_test.py -v --noconfest

===== test session starts =====
platform linux -- Python 3.12.3, pytest-9.1.1, pluggy-1.6.0 -- /home/arj/OS-Milestone-1/.venv/bin/python3
cachedir: .pytest_cache
rootdir: /home/arj/OS-Milestone-1
collecting ... collected 14 items

server/kernel_regression_test.py::test_case[thread_stack] PASSED [ 7%]
server/kernel_regression_test.py::test_case[canary] PASSED [ 14%]
server/kernel_regression_test.py::test_case[mem_rlimit_as] PASSED [ 21%]
server/kernel_regression_test.py::test_case[mem_inherit] PASSED [ 28%]
server/kernel_regression_test.py::test_case[exec_rlimit] PASSED [ 35%]
server/kernel_regression_test.py::test_case[cpu_rlimit] PASSED [ 42%]
server/kernel_regression_test.py::test_case[cpu_sleep] PASSED [ 50%]
server/kernel_regression_test.py::test_case[cap_rlimit] PASSED [ 57%]
server/kernel_regression_test.py::test_case[getentropy] PASSED [ 64%]
server/kernel_regression_test.py::test_case[wx_reject] PASSED [ 71%]
server/kernel_regression_test.py::test_case[id_priv] PASSED [ 78%]
server/kernel_regression_test.py::test_case[canary_per_exec] PASSED [ 85%]
server/kernel_regression_test.py::test_case[stack_aslr] PASSED [ 92%]
server/kernel_regression_test.py::test_case[user_tests_clean] PASSED [100%]

===== 14 passed in 253.65s (0:04:13) =====
```

Figure 6: server/kernel_regression_test.py - 14/14 passed, matching the expected table

```
pytest server/weightfetch_regression_test.py -v --noconfest

===== test session starts =====
platform linux -- Python 3.12.3, pytest-9.1.1, pluggy-1.6.0 -- /home/arj/OS-Milestone-1/.venv/bin/python3
cachedir: .pytest_cache
rootdir: /home/arj/OS-Milestone-1
collecting ... collected 4 items

server/weightfetch_regression_test.py::test_range_fetch_is_byte_exact[range_digest_aligned] PASSED [ 25%]
server/weightfetch_regression_test.py::test_range_fetch_is_byte_exact[range_digest_unaligned] PASSED [ 50%]
server/weightfetch_regression_test.py::test_undeclared_flow_is_rate_limited PASSED [ 75%]
server/weightfetch_regression_test.py::test_worker_survives_absent_master PASSED [100%]

===== 4 passed in 32.22s =====
```

Figure 7: server/weightfetch_regression_test.py - 4/4 passed, matching the expected table

```

pytest server/distinf_regression_test.py -v --noconfest

===== test session starts =====
platform linux -- Python 3.12.3, pytest-9.1.1, pluggy-1.6.0 -- /home/arj/OS-Milestone-1/.venv/bin/python3
cachedir: .pytest_cache
rootdir: /home/arj/OS-Milestone-1
collecting ... collected 8 items

server/distinf_regression_test.py::test_ring_generates FAILED [ 12%]
server/distinf_regression_test.py::test_matches_single_node FAILED [ 25%]
server/distinf_regression_test.py::test_survives_worker_loss[mid_ring] SKIPPED [ 37%]
server/distinf_regression_test.py::test_survives_worker_loss[tail] SKIPPED [ 50%]
server/distinf_regression_test.py::test_capability_lie_is_refused SKIPPED [ 62%]
server/distinf_regression_test.py::test_corrupted_output_is_not_detected[signflip] SKIPPED [ 75%]
server/distinf_regression_test.py::test_corrupted_output_is_not_detected[zero] SKIPPED [ 87%]
server/distinf_regression_test.py::test_capability_lie_above_probe_ceiling_is_accepted SKIPPED [100%]

===== FAILURES =====
_____ test_ring_generates _____

sim_output = CompletedProcess(args=['/home/arj/OS-Milestone-1/.venv/bin/python3',
'/home/arj/OS-Milestone-1/server/distinf_sim.py', ...], returncode=1, stdout='', stderr='')

    def test_ring_generates(sim_output):
        """
        The ring forms and completes a generation.

        Checks performed:
        1. The simulation reports PASS, meaning every milestone was observed in order:
           registration -> layer assignment -> shards resident -> tokens -> done.
        2. The process exits zero, so a milestone timeout fails the gate rather than
           being reported as a successful run with missing output.
        """
        > assert sim_output.returncode == 0, (
            f"simulation failed:\n{sim_output.stdout[-3000:]}\\n{sim_output.stderr[-1000:]}")
E       AssertionError: simulation failed:
E       [sim] staging 3 prebuilt node image(s)...
E       [sim] starting host mcast switch (weight server on the segment)...
E       [sim] booting 3 node(s) (timeout 240s each)...
E       [sim] master: expecting 2 worker(s), model 1 (stories15M), 32 tokens
E       [sim] OK worker1 registered (1.3s)
E       [sim] OK worker2 registered (1.3s)
E       [sim] OK 2 worker(s) registered (serialized)
E       [sim] FAIL -- never assigned layers within 900s: ['worker1', 'worker2']
E       [sim] ===== run summary =====
E       [sim] topology      master 10.0.0.2 + 2 worker(s) | model 1 (stories15M) | 32 steps
E       [sim] build          3 image(s) in 0s
E       [sim] boot           master 20s | worker1 20s | worker2 20s
E       [sim] registration   worker1 1.3s | worker2 1.3s (serialized)
E       [sim] logs           /tmp/distinf_sim/(master,worker1..N).log
E       [sim] result         FAIL -- never assigned layers within 900s: ['worker1', 'worker2']
E       [sim] =====
E       [sim] FAIL -- never assigned layers within 900s: ['worker1', 'worker2']
E
E       assert 1 == 0
E       + where 1 = CompletedProcess(args=['/home/arj/OS-Milestone-1/.venv/bin/python3',
E       '/home/arj/OS-Milestone-1/server/distinf_sim.py', '--workers', '2', '--model', '1', '--steps', '32',
E       '--prompt', 'Once upon a time'], returncode=1, stdout="[sim] staging 3 prebuilt node image(s)...\\n[sim]
E       starting host mcast switch (weight server on the segment)...\\n[sim] booting 3 node(s) (timeout 240s
E       each)...\\n[sim] master: expecting 2 worker(s), model 1 (stories15M), 32 tokens\\n[sim] OK worker1 registered
E       (1.3s)\\n[sim] OK worker2 registered (1.3s)\\n[sim] OK 2 worker(s) registered (serialized)\\n[sim] FAIL --
E       never assigned layers within 900s: ['worker1', 'worker2']\\n[sim] ===== run summary =====
E       [sim] topology      master 10.0.0.2 + 2 worker(s) | model 1 (stories15M) | 32
E       steps\\n[sim] build          3 image(s) in 0s\\n[sim] boot           master 20s | worker1 20s | worker2
E       20s\\n[sim] registration   worker1 1.3s | worker2 1.3s (serialized)\\n[sim] logs           /tmp/distinf_sim/(master,worker1..N).log
E       [sim] result         FAIL -- never assigned layers within 900s: ['worker1', 'worker2']\\n[sim] =====
E       [sim] FAIL -- never assigned layers within 900s: ['worker1', 'worker2']\\n", stderr='').returncode
E
server/distinf_regression_test.py:95: AssertionError
_____ test_matches_single_node _____

sim_output = CompletedProcess(args=['/home/arj/OS-Milestone-1/.venv/bin/python3',
'/home/arj/OS-Milestone-1/server/distinf_sim.py', ...], returncode=1, stdout='', stderr='')

    def test_matches_single_node(sim_output):
        """
        The distributed token stream equals the single-node golden reference.

        Checks performed:
        1. The generated text is recoverable from the master's console.
        2. It matches the reference exactly -- not approximately. A mismatch means the
           sharded path diverged from the whole-model path; wrong layer range, wrong KV
           cache indexing, or an activation corrupted in transit.
        """
        > generated = generated_text(sim_output.stdout)
E
server/distinf_regression_test.py:134:
-----
out = "[sim] staging 3 prebuilt node image(s)...\\n[sim] starting host mcast switch (weight server on the
segment)...\\n[sim] .....\\n[sim] FAIL -- never assigned layers within
900s: ['worker1', 'worker2']\\n"

    def generated_text(out):
        """The token block the simulation printed between its two banner lines."""
        start = out.find("==== generated =====")
        end = out.find("====")
        > assert start >= 0 and end > start, f"no generated text in:\\n{out[-2000:]}"
E       AssertionError: no generated text in:
E       [sim] staging 3 prebuilt node image(s)...
E       [sim] starting host mcast switch (weight server on the segment)...
E       [sim] booting 3 node(s) (timeout 240s each)...
E       [sim] master: expecting 2 worker(s), model 1 (stories15M), 32 tokens
E       [sim] OK worker1 registered (1.3s)
E       [sim] OK worker2 registered (1.3s)
E       [sim] OK 2 worker(s) registered (serialized)
E       [sim] FAIL -- never assigned layers within 900s: ['worker1', 'worker2']
E       [sim] ===== run summary =====
E       [sim] topology      master 10.0.0.2 + 2 worker(s) | model 1 (stories15M) | 32 steps
E       [sim] build          3 image(s) in 0s
E       [sim] boot           master 20s | worker1 20s | worker2 20s
E       [sim] registration   worker1 1.3s | worker2 1.3s (serialized)
E       [sim] logs           /tmp/distinf_sim/(master,worker1..N).log
E       [sim] result         FAIL -- never assigned layers within 900s: ['worker1', 'worker2']
E       [sim] =====
E       [sim] FAIL -- never assigned layers within 900s: ['worker1', 'worker2']
E
E       assert (-1) >= 0
E
server/distinf_regression_test.py:104: AssertionError
===== short test summary info =====
FAILED server/distinf_regression_test.py::test_ring_generates - AssertionError...
FAILED server/distinf_regression_test.py::test_matches_single_node - Assertio...
===== 2 failed, 6 skipped in 523.38s (0:08:43) =====

```

Figure 8: server/distinf_regression_test.py – failed reproducibly (3/3 attempts) with the same environment-level multicast flakiness; see Observations Log #1 and #5 and Contributions #4. A manual distinf_sim.py invocation with identical arguments succeeded and reproduced the single-node token stream exactly (see Bring-Up Evidence above).

Architecture Map (Section 3.3) — Eight Subsystems

Method: rows 2-7 are read directly from user/*.c source (file:line below). Rows 1 and 8 live in the kernel, whose .c source is withheld, so they are mapped from the binary — riscv64-unknown-elf-nm kernel/kernel for addresses and riscv64-unknown-elf-objdump -d --disassemble=<fn> kernel/kernel for behaviour.

1. Network Stack (kernel, from binary)

- **Entry point:** net_rx, address 0x800099c6 (kernel.sym)
- **Responsibility:** dispatches every received Ethernet II frame by EtherType.
- **From the disassembly:** reads bytes 12-13 of the frame (the EtherType field) as a 16-bit value. 0x0806 (ARP) branches to arp_rx (0x800094e8); 0x0800 (IPv4) branches to ip_rx (0x80008d9a); anything else is dropped via kfree.
- **Standard:** Ethernet II framing; the two dispatch targets implement RFC 826 (ARP) and RFC 791 (IPv4), per the header comments in kernel/net.h.

2. Remote Procedure Call

- **Entry point:** rpc_call, user/rpc.c:487
- **Responsibility:** ONC RPC transport over UDP — request/reply framing, retry.
- **Standard:** RFC 5531 (ONC RPC), quoted directly in the file header comment (user/rpc.c:1-6): “ONC RPC transport layer, RFC 5531 compliant... Record marking (RFC 5531 §11) is TCP-only; not implemented here” (this stack is UDP-only).

3. Data Serialisation

- **Entry point:** xdrmem_create, user/xdr.c:143
- **Responsibility:** in-memory XDR encode/decode stream backing every RPC message.
- **Standard:** RFC 4506 (External Data Representation), referenced in user/rpc.c’s header comment (“XDR encoding per RFC 4506”).

4. Discovery and Registration

- **Entry point:** discovery_register, user/discovery.c:288
- **Responsibility:** worker→master registration handshake; proves PSK possession via HMAC before a worker is admitted to the ring.
- **Standard:** RFC 2104 / FIPS 198-1 (HMAC), quoted directly at user/discovery.c:8: “hmac_sha256 — RFC 2104 / FIPS 198-1 keyed MAC over user/sha256.c.”

5. Ring Driver and Shard Logic

- **Entry points:** run_master, user/distinf.c:1268; run_worker, user/distinf.c:1392; shard_load, user/shard.c:103; shard_forward, user/shard.c:190.
- **Responsibility:** distinf.c drives registration, layer assignment, and the token-by-token ring traversal; shard.c computes each worker’s layer range within a checkpoint and loads only that slice.

6. Transformer Kernels

- **Entry point:** forward, user/llama_core.c:742
- **Responsibility:** the transformer forward pass (attention, matmul, rmsnorm) shared by both the single-node driver (llama.c) and each ring shard.

7. Weight Fetch Client

- **Entry point:** `llm_fetch_file`, `user/ftpclient.c:691`
- **Responsibility:** fetches a checkpoint/tokenizer file over the custom UDP file-transfer protocol (LLM-RFTP) from `server/server.py`, chunk by chunk, verifying the transfer against a SHA-256 the server reports at connect time.

8. Kernel Hardening Set (kernel, from binary)

Feature	Entry point	Address	What the disassembly shows
Stack canary init	<code>stackguard_init</code>	<code>0x80001e64</code>	Calls <code>rand64</code> (<code>0x80001f3c</code>), stores the result to <code>__stack_chk_guard</code> (<code>0x8000cbb0</code>) — a real per-boot random canary, not a constant.
Canary check	<code>__stack_chk_fail</code>	<code>0x80001e80</code>	Called from guarded functions (e.g. <code>sys_exec</code> at <code>0x80005958</code>) when the saved canary no longer matches.
Entropy pool init	<code>random_init</code>	<code>0x80002000</code>	Seeds the pool by calling <code>random_add</code> 8 times, each time mixing in <code>rdtime</code> (the RISC-V cycle/time CSR) plus stack contents — a real boot-time entropy source, not a fixed seed.
User-facing entropy syscall	<code>sys_getentropy</code>	<code>0x80001b46</code>	Standard syscall preamble (<code>argaddr/argint</code> argument marshalling) around the kernel's entropy pool — the user-side half of the canary/key-randomisation story the README describes.

Feature	Entry point	Address	What the disassembly shows
Memory quota (RLIMIT_AS)	rlimit_as_ok	0x8000693a	Compares a process struct field at offset +72 (current usage) against the requested size, and a field at +656 (the limit) — returns 1 only if the request stays under the hard limit.
Capabilities	capable	0x800069d6	Loads the calling process's capability bitmask (proc struct offset +696), shifts right by the requested bit index, masks bit 0 — single-bitmask capability check, matching the README's "a single effective bitmask suffices for the one capability in use."
W^X enforcement	kexec	0x8000489e	At 0x80004a96: masks a segment's permission bits (flags2perm result) with 0xc (PTE_W \ PTE_X from kernel/riscv.h:410-411) and compares to 0xc — a segment requesting both writable and executable is rejected before it is ever mapped.

Resource-limit syscalls (sys_setrlimit/sys_getrlimit, 0x80001982/0x80001a44) and sys_exec (0x80005838, which calls into kexec for the actual ELF load and W^X check) round out the set; all addresses above are directly from riscv64-unknown-elf-nm kernel/kernel.

Figures (Section 3.4)

F1 — System Architecture (flowchart)

Weight server, master, the N workers, the ring, and the transport between them, drawn from `server/server.py`, `server/mcast_switch.py/boot.sh`'s SLIRP networking, and `user/distinf.c`'s master/worker roles.

flowchart LR

```
    subgraph HOST["Host (Linux / server/)"]
        WS["Weight server<br>server/server.py<br>UDP :9999, LLM-RFTP"]
        SW["mcast_switch.py<br>(ring config only: bridges WS<br>onto the QEMU
        ↪ multicast segment,<br>answers ARP for 10.0.0.1)"]
    end

    subgraph GUEST0["xv6 node: master"]
        MA["distinf --master<br>run_master() :1268<br>registry, layer
        ↪ assignment,<br>master_run_token per-token loop"]
        LC0["llama_core.c<br>embedding table + classifier only"]
    end

    subgraph GUEST1["xv6 node: worker 1 (head)"]
        W1["distinf --worker<br>run_worker() :1392"]
        SH1["shard.c: layers [0,k)"]
        LC1["llama_core.c: llama_block per owned layer"]
    end

    subgraph GUESTN["xv6 node: worker N (tail)"]
        WN["distinf --worker"]
        SHN["shard.c: layers [j,n_layers)"]
        LCN["llama_core.c: llama_block per owned layer"]
    end

    WS -- "LLM-RFTP: tokenizer.bin, stories*.bin<br>user/ftpcient.c" --> MA
    WS -- LLM-RFTP --> W1
    WS -- LLM-RFTP --> WN
    SW -. "bridges WS onto mcast segment<br>(multi-worker ring only)" .-> WS

    MA == "PROC_INFER_REQ (ring hop 0)<br>RPC/UDP, RFC 5531" ==> W1
    W1 == "PROC_INFER_REQ (hop 1)" ==> WN
    WN == "PROC_INFER_REQ (hop N, INFER_FLAG_FINAL)" ==> MA

    MA -. "PROC_AUTH_HELLO / PROC_CAP_ACK<br>discovery.c registration" .-> W1
    MA -. registration .-> WN
    MA -. "PROC_ASSIGN_LAYERS" .-> W1
    MA -. "PROC_ASSIGN_LAYERS" .-> WN
    W1 -. "PROC_SHARD_READY" .-> MA
    WN -. "PROC_SHARD_READY" .-> MA
```

Solid double arrows: the per-token ring hop path (the hot path, F4/F5). Dashed arrows: control-plane RPCs (registration, layer assignment, shard-ready) that run once per session, not once per token. The single-node configuration (`llama.c`) is the degenerate case of this diagram with zero workers: `llama_core.c` runs every layer locally and only the weight-server edge exists.




Figure 9: F1: System architecture

F2 — Wire Format: RPC Envelope + XDR Activation Message (packet-beta)

Source: `user/rpc.h` (RFC 5531 §9 `call_body/rpc_msg` envelope) and `user/distinf.h` (`infer_hop_t`, the `PROC_INFER_REQ` payload — one ring-hop activation message). All multi-byte fields are 32-bit big-endian per RFC 4506 §4.2; the activation array is IEEE 754 binary32, canonical big-endian, RFC 4506 §4.6/§4.12.

packet-beta

title `PROC_INFER_REQ` on the wire: RPC envelope (RFC 5531 §9) + XDR activation

↪ payload (RFC 4506)

```

0-31: "xid (transaction id)"
32-63: "mtype = CALL (0)"
64-95: "rpcvers = 2"
96-127: "prog = INFERENCE_PROG"
128-159: "vers = INFERENCE_VERS"
160-191: "proc = PROC_INFER_REQ (4)"
192-223: "cred.flavor = AUTH_NONE (0)"
224-255: "cred.body_len = 0"
256-287: "verf.flavor = AUTH_NONE (0)"
288-319: "verf.body_len = 0"
320-351: "session_id (epoch)"
352-383: "seq (strictly increasing, anti-replay)"
384-415: "pos (token position, RoPE + KV index)"
416-447: "hop (expected handler index)"
448-479: "n_floats (= dim)"
480-511: "flags (bit0 = INFER_FLAG_FINAL)"
512-543: "compute_us (cumulative shard_forward time)"
544-575: "activation[0] (float32, big-endian, RFC 4506 §4.6/§4.12)"
576-607: "activation[1] (float32, big-endian)"
608-639: "..."

```

The envelope (`xid` through `verf.body_len`, 40 bytes) is identical across every RPC procedure; only `proc` (word 5) and the payload after it change. The payload here (`infer_hop_t`, `distinf.h`:87-95, 7 words = 28 bytes) is followed by `n_floats` activation floats — at `dim`=768 that is 3072 bytes of floats, for a total payload of 3100 bytes, deliberately kept

0	xid (transaction id)	31
32	mtype = CALL (0)	63
64	rpcvers = 2	95
96	prog = INFERENCE_PROG	127
128	vers = INFERENCE_VERS	159
160	proc = PROC_INFER_REQ (4)	191
192	cred.flavor = AUTH_NONE (0)	223
224	cred.body_len = 0	255
256	verf.flavor = AUTH_NONE (0)	287
288	verf.body_len = 0	319
320	session_id (epoch)	351
352	seq (strictly increasing, anti-replay)	383
384	pos (token position, RoPE + KV index)	415
416	hop (expected handler index)	447
448	n_floats (= dim)	479
480	flags (bit0 = INFER_FLAG_FINAL)	511
512	compute_us (cumulative shard_forward time)	543
544	activation[0] (float32, big-endian, RFC 4506 §4.6/§4.12)	575
576	activation[1] (float32, big-endian)	607
608	...	639

PROC_INFER_REQ on the wire: RPC envelope (RFC 5531 §9) + XDR activation payload (RFC 4506)

Figure 10: F2: PROC_INFER_REQ wire format

under `RPC_PAYLOAD_MAX` (4096, `rpc.h:177`) but above the 1500-byte Ethernet MTU, so the header comment in `distinf.h:22-25` notes this is by design: every hop exercises IP fragmentation/reassembly (RFC 791) rather than leaving that path untested.

F3 — Worker Lifecycle (stateDiagram-v2)

Built from `user/discovery.h` (states, constants) and `user/discovery.c` (transitions), not from `diagrams/state_diagram.png`, which does not match the code (see divergence paragraph below).

Real timing constants (`user/discovery.h:18-25`): `HEARTBEAT_INTERVAL_TICKS = 200`, `SUSPECTED_TIMEOUT = 400` ticks (2 missed heartbeats), `EXPIRED_TIMEOUT = 800` ticks (4 missed heartbeats). These are **ticks**, not milliseconds, despite every call site passing `uptime()` into a parameter named `now_ms` (see Observations Log #2) — the diagram below labels transitions with tick counts to stay accurate to the code.

stateDiagram-v2

```
[*] --> PENDING: AUTH_HELLO, valid PSK commitment (discovery.c:504-506)
PENDING --> ACTIVE: CAP_ACK valid (nonce + RAM-probe checksum + HMAC identity
↪ all match, discovery.c:580-581)
PENDING --> SUSPECTED: age > 400 ticks, no valid CAP_ACK yet
↪ (discovery.c:795-798)
ACTIVE --> SUSPECTED: age > 400 ticks since last heartbeat
↪ (discovery.c:795-798)
SUSPECTED --> ACTIVE: heartbeat received (discovery.c:660-662)
PENDING --> EXPIRED: age > 800 ticks (discovery.c:784-786)
SUSPECTED --> EXPIRED: age > 800 ticks (discovery.c:784-786)
EXPIRED --> PENDING: worker re-sends AUTH_HELLO, slot reused
↪ (discovery.c:482,494-496)
EXPIRED --> EXPIRED: heartbeat received while EXPIRED -- rejected, SYSTEM_ERR
↪ (discovery.c:664-666)
```

Where `diagrams/state_diagram.png` is wrong

The shipped hand-drawn diagram invents a fifth state, “unregistered worker,” with a worker pending --timeout--> unregistered worker edge (the slot is freed if the capability probe is never completed) and a separate hello edge back into pending. Neither exists in the registry code:

1. **There is no “unregistered” state.** An unregistered worker simply has no registry entry (`registry[i].info.worker_id == 0` is the sentinel for an empty slot, not a tracked state) — `WORKER_PENDING/ACTIVE/SUSPECTED/EXPIRED` are the only four values `worker_entry_t.state` ever takes (`discovery.h:13-16`).
2. **A PENDING worker that never completes the capability probe is not freed on timeout.** `discovery_expire_workers` (`discovery.c:769-800`) does not special-case `WORKER_PENDING` at all — it ages *every* non-`EXPIRED` slot by the same `SUSPECTED_TIMEOUT/EXPIRED_TIMEOUT` thresholds regardless of current state, so a stalled `PENDING` worker is walked through `PENDING → SUSPECTED → EXPIRED` exactly like a stalled `ACTIVE` one, not dropped back to “unregistered.” This is not a guess: the code says so directly, in a comment at `discovery.c:777-782`, which explicitly calls this out as “a known divergence from the maintainer’s lifecycle diagram” — i.e. the bug in the diagram (not the code) was already known and left undocumented anywhere but that one inline comment.
3. **Re-registration is EXPIRED → PENDING, not “→ unregistered → pending.”** When an `EXPIRED` worker sends a fresh `AUTH_HELLO`, `handle_auth_hello` finds its existing (expired) slot and reuses it directly, setting the state straight back to `PENDING` (`discovery.c:482, 494-496, 506`) — there is no intermediate state.




Figure 11: F3: Worker lifecycle state machine

F4 — One Token, Master → Worker1 → Worker2 → Master (sequenceDiagram)

Shown for a 2-worker ring (matches the captured bring-up transcript), real function names and RPC procedures from user/distinf.c.

sequenceDiagram

```

    participant M as Master (master_run_token, :1067)
    participant W1 as Worker 1 (head, layers [0,3))
    participant W2 as Worker 2 (tail, layers [3,6))

    M->>M: llama_embed(x, token) :1072
    M->>W1: PROC_INFER_REQ infer_send_hop(head, h, x) :1102
    W1->>W1: shard_forward(shard, compute, pos) :519
    W1->>W2: PROC_INFER_REQ infer_send_hop(next, out, compute) :569
    W2->>W2: shard_forward(shard, compute, pos) :519
    W2->>M: PROC_INFER_REQ infer_send_hop(master, out, compute) :569
    M->>M: master_pump(200) receives it, sets g_have_final :1109-1111
    M->>M: g_metric_rtt_us/compute_us/tokens updated, return 0 :1114-1117

```




Figure 12: F4: One token, master through each worker and back

Retry note (not shown per-hop above): if no reply reaches the master within `INFER_TOKEN_TIMEOUT_MS` (15000 ms, `distinf.h:46`), `master_run_token` retries the *same* token up to `INFER_MAX_RETRIES` (2) times, each attempt with a freshly incremented `seq` so a worker's anti-replay check accepts the retransmit rather than rejecting it as a replay (`distinf.c:1080-1099`) — the ring itself is fire-and-forget (RFC 5531 §5 batching); timeout/retry is the master's job alone.

F5 — Inference Hot Path Through user/llama_core.c (flowchart)

Call graph rooted at forward() (llama_core.c:742), in the style of syscall_dependency_diagram.md. Every edge is a real call in the source; matmul and attention each have two paths (serial fallback vs. thread-pool dispatch), both shown.

graph TD

```
F["forward(transformer, token, pos) llama_core.c:742"] -->
  ↳ EMB["llama_embed(w,p,x,token) :640"] --> MEMCPY["memcpy(x, embedding_row)"]
F --> LA["llama_layer_at(w,p,l) :619 (pointer arithmetic only, per layer)"]
F --> BLK["llama_block(layer,p,s,x,l,pos) :653 (per layer, 0..n_layers)"]

BLK --> RN1["rmsnorm(xb,x,rms_att,dim) :662 (attention rmsnorm)"]
BLK --> MMQ["matmul(q,xb,wq,dim,dim) :670"]
BLK --> MMK["matmul(k,xb,wk,dim,kv_dim) :671"]
BLK --> MMV["matmul(v,xb,wv,dim,kv_dim) :672"]
BLK --> MHA["multihead_attention(s,p,0,l,pos) :694"]
BLK --> MM0["matmul(xb2,xb,wo,dim,dim) :697"]
BLK --> RN2["rmsnorm(xb,x,rms_ffn,dim) :705 (ffn rmsnorm)"]
BLK --> MM1["matmul(hb,xb,w1,dim,hidden_dim) :709"]
BLK --> MM3["matmul(hb2,xb,w3,dim,hidden_dim) :710"]
BLK --> MM2["matmul(xb,hb,w2,hidden_dim,dim) :723"]

MMQ --> MMFORK{"d < 128 or pool uninitialised? :286"}
MMFORK -- yes: serial --> MMSE["inline dot-product loop :287-294"]
MMFORK -- no: parallel --> MMDISP["dispatch TASK_MATMUL to thread pool :307-329"]
  ↳ --> UWT["universal_worker_thread :160"] -->
  ↳ ↳ DPU1["dot_product_unrolled(w_row,x,n) :175, :81"]

MHA --> MHAFORK{"pos < 32 or pool uninitialised? :568"}
MHAFORK -- yes: serial --> WDA1["worker_do_attention(seq_work) :576, :110"]
MHAFORK -- no: parallel --> ATTDISP["dispatch TASK_ATTENTION to thread pool
  ↳ :587-605"] --> UWT --> WDA2["worker_do_attention(att_work) :179, :110"]

WDA1 --> DPU2["dot_product_unrolled(q,k,head_size) :132"]
WDA1 --> SM1["softmax(att,pos+1) :139, :533"]
WDA2 --> DPU2
WDA2 --> SM1

F --> HEAD["llama_head(w,p,s,x) :732"]
HEAD --> RN3["rmsnorm(x,x,rms_final,dim) :735"]
HEAD --> MMCLS["matmul(logits,x,wcls,dim,vocab_size) :738"]
```

Notes: UWT (the thread-pool worker loop, :160) is drawn once and reached from both dispatch sites — it is the function called from two different places in this graph, mirroring the mechanism the warm-up exercise asked to identify in miniature. dot_product_unrolled (:81) is the innermost kernel, reached on every code path (serial matmul, pooled matmul, and both attention paths).




Figure 13: F5: Inference hot path through llama_core.c

F6 — Node Memory Map (block-beta)

Physical layout from kernel/memlayout.h; virtual layout top addresses (MAXVA, TRAMPOLINE, TRAPFRAME) from kernel/riscv.h:427-430 and memlayout.h:65-67; the trampoline-region placement is independently confirmed live — the gdb trace captured for the reverse-engineering task (docs/transcripts/gdb_kernel_trace.txt) shows frame #3 returning to 0x3fffffff124, inside TRAMPOLINE = MAXVA - PGSIZE = 0x3fffffff000, exactly where the static header says user code re-enters the kernel after a trap.

Two address spaces, drawn as two separate columns (physical RAM the kernel manages directly; the per-process virtual address space every user program, including llama, runs in) rather than one continuous range, since a physical address and a virtual address here are never the same number.

block-beta

columns 2

```
PHYS_HDR["Physical RAM"] VIRT_HDR["Process virtual address space (high to low)"]
PT["PHYSTOP = KERNBASE + 256 MiB\n= 0x90000000\n(RAM_IN_MB=256,
↪ memlayout.h:44)"] TR["TRAMPOLINE = MAXVA-PGSIZE\n= 0x3fffffff000\n(shared
↪ kernel/user trap entry page)"]
KHEAP["kernel page-allocation\narea (kalloc pool)"] TF["TRAPFRAME =
↪ TRAMPOLINE-PGSIZE\n= 0x3ffffe000\n(p->trapframe)"]
KTEXT["kernel text + data\n(from 'end' symbol\ndown to here)"] STK["user
↪ stack(s): main thread +\nworker-pool thread stacks.\nntthread_create shares user
↪ memory\nwith the main process, so pool-thread\nstacks are carved from the
↪ same\ngrowable region, not separate\naddress spaces"]
KB["KERNBASE = 0x80000000\nentry.S, kernel image\nloaded here"] HEAP["heap,
↪ grows up via sbrk (growproc).\nModel weights (TransformerWeights,\nnmmap'd from
↪ shared-memory cache) +\nRunState activations
↪ (malloc_run_state,\nllama_core.c:339) live here. Bounded by\nRLIMIT_AS
↪ (rlimit.h:30) -- growproc fails\nonce mapped size would exceed the
↪ soft\nlimit, independent of physical RAM"]
space BSS["bss + original data"]
space TEXT["text (program code)"]
space ZERO["address 0"]
```

Only the physical-layout block and the two fixed top-of-VA pages (TRAMPOLINE, TRAPFRAME) are constants read from the kernel headers/binary; the heap/stack split below them is the ordinary xv6 user layout (memlayout.h's own comment block, "text / data+bss / fixed-size stack / expandable heap") with the model weights and activation buffers identified as the heap's largest occupants from malloc_run_state/memory_map_weights in user/llama_core.c. RLIMIT_AS is



Figure 14: F6: Node memory map, physical and virtual

the address-space limit this figure is captioned against: it caps how far HEAP may grow regardless of how much physical RAM the host QEMU process is given.

F7 — Baseline Graphs (xychart-beta)

Plots of the measurements in docs/figures/baseline_measurements.md.

Single-node inference time per prompt (ms)

xychart-beta

```
title "Single-node inference time by prompt (XV6_CPUS=1)"
x-axis ["Once upon a time", "The quick brown fox", "In the beginning", "Mean"]
y-axis "Inference time (ms)" 0 --> 14000
bar [8624, 11770, 7266, 9220]
```




Figure 15: F7a: Single-node inference time by prompt

Ring generation time: single-node mean vs. 2-worker vs. 3-worker (ms)

xychart-beta

```
title "Generation time for 32 tokens, first prompt: single-node vs. ring"
x-axis ["single-node (mean)", "2-worker ring", "3-worker ring"]
y-axis "Time (ms)" 0 --> 100000
bar [9220, 79200, 93300]
```




Figure 16: F7b: Generation time, single-node vs. ring

Per-token latency breakdown, ring only: compute vs. network ($\mu\text{s}/\text{token}$)

xychart-beta

title "Per-token latency breakdown: compute vs. network wire time"

x-axis ["2-worker ring", "3-worker ring"]

y-axis "Microseconds per token" 0 --> 2200000

bar "compute_us" [466767, 382276]

bar "network (rtt - compute)" [1638610, 2123684]




Figure 17: F7c: Per-token compute vs. network breakdown

Reading the three together: single-node is fastest in absolute terms at this model scale (F7.2); adding workers makes each individual token slower, not faster, because per-token network round-trip time on the emulated segment (F7.3) grows faster with worker count than the compute time saved by sharding shrinks — pipeline parallelism here is bottlenecked by transport, not compute, exactly as the Method section of the baseline measurements notes.

Difference From Upstream xv6 (Section 3.5)

1. Added System Calls

Stock xv6-riscv defines syscalls 1-21 (SYS_fork ... SYS_close, kernel/syscall.h:2-22). This build goes to 59, so **38 syscalls (numbers 22-59) were added**, matching the handout's estimate exactly. Numbers 22-28 (trace, interpose, sigalarm, sigreturn, symlink, mmap, munmap) have **no corresponding stub in user/user.h** — they are reserved numbers inherited from the upstream MIT 6.828 xv6-labs skeleton this project descends from, but this project never wired a userland entry point to them, so they are listed for completeness but are not callable and not part of this project's own work.

#	Name	Description
22	trace	<i>(no user-space stub; unused xv6-labs leftover)</i>
23	interpose	<i>(no user-space stub; unused xv6-labs leftover)</i>
24	sigalarm	<i>(no user-space stub; unused xv6-labs leftover)</i>
25	sigreturn	<i>(no user-space stub; unused xv6-labs leftover)</i>
26	symlink	<i>(no user-space stub; unused xv6-labs leftover)</i>
27	mmap	<i>(no user-space stub; unused xv6-labs leftover)</i>
28	munmap	<i>(no user-space stub; unused xv6-labs leftover)</i>
29	bind	Bind the calling process to a UDP port for raw send/recv.
30	unbind	Release a previously bound UDP port.
31	send	Send a UDP datagram (port, dest IP, dest port, buffer, length).
32	recv	Receive a UDP datagram on a bound port (blocking).
33	pgpte	Return the PTE for a given virtual address (page-table introspection).
34	kpgtbl	Print the kernel page table (debug/introspection).
35	shmget	Create/look up a named shared-memory segment (IPC_CREAT/IPC_EXCL).
36	shmat	Attach a shared-memory segment into the caller's address space.
37	shmdt	Detach a shared-memory segment.
38	shmctl	Control a shared-memory segment (e.g. IPC_RMID).

#	Name	Description
39	rdcycle	Read the RISC-V cycle CSR (baseline measurement primitive).
40	rdtime	Read the RISC-V time CSR.
41	rdinstret	Read the RISC-V instret CSR (instructions retired).
42	getramused	Report the caller's current RAM usage in bytes.
43	setrlimit	Set a POSIX resource limit (RLIMIT_AS/RLIMIT_CPU) for the caller.
44	getrlimit	Read a POSIX resource limit for the caller.
45	capget	Return the caller's capability bitmask.
46	capdrop	Permanently clear one capability bit for the caller (one-way).
47	getentropy	Fill a buffer with real boot-time entropy (canaries, ASLR, keys).
48	getuid	Return the caller's real uid.
49	geteuid	Return the caller's effective uid.
50	setuid	Set the caller's real (and effective) uid.
51	seteuid	Set the caller's effective uid only.
52	recvtimeo	Receive a UDP datagram with a timeout (bounded blocking recv).
53	ip	Return the caller node's own IPv4 address.
54	thread_create	Spawn a new thread sharing the caller's address space.
55	thread_join	Block until the given thread exits.
56	thread_exit	Terminate the calling thread.
57	yield	Voluntarily give up the CPU (used by the matmul/attention thread pool's spin-wait).
58	setpriority	Set a process's scheduling priority.
59	udp_bulk	Declare a bulk UDP flow (port + peer IP) exempt from the per-source rate limiter; privileged (CAP_NET_ADMIN).

A note from the kernel header itself (kernel/syscall.h:66-74, quoted verbatim): the thread syscalls (54-57) were originally numbered 43-47, directly colliding with the

setrlimit/getrlimit/capget/capdrop/getentropy block above. Because syscalls[] is a designated-initializer array, the collision was not a compile error — the later initializer silently won, so all five of those POSIX syscalls dispatched to thread/scheduler handlers instead of their intended implementations until the thread block was renumbered to 54-57. Already fixed in this build (the numbering above is current), but worth noting as a real historical defect class (silent designated-initializer collision) the maintainers document rather than hide.

2. File Classification

Grouping the added user/*.c files (and the kernel hardening set, mapped from the binary in Section 3.3) by what they are for:

Category	Files
Networking	kernel/net.c*, user/nettest.c, user/netecho.c, user/rpc.c/rpc.h, user/xdr.c/xdr.h
Distributed inference	user/discovery.c/discovery.h, user/distinf.c/distinf.h, user/shard.c/shard.h, user/dorphan.c, user/forphan.c
Model computation	user/llama_core.c/llama_core.h, user/llama.c, user/ftpclient.c/ftpclient.h, user/shardspike.c
Hardening	kernel/stackguard.c, kernel/random.c, kernel/rlimit.h, kernel/capability.h, user/canarytest.c, user/aslrtest.c, user/rlimittest.c, user/idtest.c, user/wxtest.c/wxbad.c, user/stackguard.c, user/randtest.c
Measurement	user/perf.c/perf.h, user/testperf.c, user/testram.c, user/testshm.c
Test/regression infra (cross-cutting)	user/testutil.c/testutil.h, user/runtests.c, user/usertests.c, user/rpc_test.c, user/xdr_test.c, user/testftp.c, user/testsha.c, user/mutex.c, user/sha256.c/sha256.h, user/xmath.c/xstdlib.c/xstrlib.c (freestanding math/libc replacements, needed because the kernel/user split here has no libc)

* kernel-side file; .c source withheld, classified from its role as established in Section 3.3 (binary-mapped rows) rather than by reading it directly.

Baseline Measurements (Section 3.6)

Method

Fixed workload throughout, per the handout: model stories15M.bin (id 1), steps=32, temperature=0 (-t 0, always). Single-node runs at XV6_CPUS=1 (boot.sh defaults to 3 cores; every measured single-node run here explicitly passed XV6_CPUS=1 ./boot.sh). Ring runs use distinf_sim.py's default --smp 8 per node (unchanged from default so the 2- and 3-worker figures are measured under the same per-node core budget as each other).

Primary number: wall-clock duration, taken directly from each run's own reported timing — Inference (ms) from llama's built-in profiler for single-node, and the generation ... in Xs line from distinf_sim.py's run summary for the ring. Using Inference rather than Total (E2EL) for single-node excludes the one-time network weight-fetch cost (dominated by emulated-UDP transfer time, not the model's compute cost), matching what the ring's generation time also measures (compute + ring network hops for the shard path, excluding the separate shard *loading* phase that precedes it).

Secondary numbers: cycles and instructions retired, computed rather than sampled live per run. This needs justification: rdcycle/rdinstret were tested directly with gdb against the halted kernel (docs/transcripts/cycle_instret_probe.txt, Observations Log #5) and found to advance in lockstep with real elapsed time *even while the guest vCPU was provably halted and executed zero instructions* — under this QEMU configuration (no -icount), these CSRs are wall-clock-derived, not true per-instruction counters, and empirically instret $\approx$ cycle throughout (implied IPC $\approx$ 1.00 at all times, including idle). A live per-run gdb sample would therefore carry no information beyond the wall-clock time already measured. Cycles and instructions retired are instead reported as wall_clock_seconds $\times$ 2.2597×10^9 , using the scale factor established by that same probe (two samples, 5.00s apart, Δ instret = 11,298,623,360, Δ cycle = 11,298,625,329). This is reported per the handout's own instruction to favour ratios/relative comparisons over absolute values — the derived figures are internally consistent for that purpose even though they are not literal hardware-counted values.

Single-Node (XV6_CPUS=1)

Prompt	Inference (ms)	Cycles (derived)	Instructions retired (derived)
"Once upon a time"	8,624	1.949×10^{10}	1.949×10^{10}
"The quick brown fox"	11,770	2.660×10^{10}	2.660×10^{10}
"In the beginning"	7,266	1.642×10^{10}	1.642×10^{10}
Mean	9,220	2.083×10^{10}	2.083×10^{10}

Source: docs/transcripts/baseline_singlenode.txt (all three prompts run in one boot session, weights fetched once and cached — see the Inference field of each run's own SYSTEM PERFORMANCE METRICS block; Total (E2EL) for the first prompt additionally includes the one-time ~76.6s weight fetch, excluded here as explained in Method).

Run-to-run variance is real and worth flagging: identical workload (32 steps, same model, same 1,376 matmul calls per run per the function-profiling table), yet per-run matmul self-time ranged 7,085–11,458 ms across the three prompts — a ~40% swing attributable to host scheduling noise (this machine is a shared, CPU- oversubscribed WSL2 host throughout this

session; see Observations Log #1/#5), not to the workload itself. The mean is reported precisely because of this variance.

Multi-Node Ring (first prompt, “Once upon a time”)

Configuration	Generation time (32 tokens)	Cycles (derived)	Instructions retired (derived)	Compute-only (sum of com- pute_us/token × 32)	Network (sum of net- work/token × 32)
2 workers	79.2 s	1.790×10^{11}	1.790×10^{11}	14.94 s	52.44 s
3 workers	93.3 s	2.108×10^{11}	2.108×10^{11}	12.23 s	67.96 s

Source: docs/transcripts/ring_2workers_transcript.txt, docs/transcripts/baseline_ring_3workers.tx
— each run’s own run summary block (generation, latency/token (avg over 32): ring rtt
... = compute ... + network ...). Both rings reproduced the single-node token stream
exactly (“Once upon a time, there was a little girl named Lily...”).

The ring is slower in wall-clock terms than the single-node mean, and 3 workers is slower than 2 — the opposite of what “more parallelism” might suggest, but expected here: each ring hop pays a full RPC round trip over an emulated network segment (network time dominates per-token latency in both configurations — 66% at 2 workers, 73% at 3 — and grows with worker count as expected, since more workers means more hops per token), while compute time per token *drops* as workers increase (each worker does less of the model: w1 [0,2) | w2 [2,4) | w3 [4,6) at 3 workers vs. w1 [0,3) | w2 [3,6) at 2), the opposite direction from overall latency. This is the expected shape for pipeline parallelism at small model scale over a slow transport: it only pays off once per-shard compute exceeds the added network cost, which does not hold for a 15M-parameter model on an emulated UDP segment. The measurement directly supports the project’s own stated v1 scope note (sequential per-token pipeline, no batching) — throughput, not latency, is where pipeline parallelism would need to be evaluated for this to look favourable, and that is out of scope for a single generation run.

```
gdb: cycle/instret CSRs sampled while vCPU genuinely halted

$ gdb-multiarch -batch -x probe.gdb xv6-riscv/kernel/kernel
# probe.gdb: target remote localhost:26000 ; info registers instret cycle mcycle minstret time
# CPU is halted (-S), never sent 'continue' -- zero guest instructions executed.

--- Sample 1 (t=0s, immediately after connecting) ---
instret      0x17cdb114779  1635762915193
cycle        0x17cdb1d753f  1635763713343
mcycle       0x17cdb26e502  1635764331778
minstret     0x17cdb31c928  1635765045544
time         0x0           0

--- (5 real seconds elapse; CPU remains halted the entire time) ---

--- Sample 2 (t=5s) ---
instret      0x17f7c8496f9  1647061538553
cycle        0x17f7c90cc70  1647062338672
mcycle       0x17f7c9c63d4  1647063098324
minstret     0x17f7ca72ce8  1647063805160
time         0x30d503a     51204154

--- Delta over 5.00s of real time, 0 guest instructions executed ---
instret: +11298623360 (~2.2597 GHz)
cycle:   +11298625329 (~2.2597 GHz)
```

Figure 18: gdb probe: cycle/instret CSRs sampled while the vCPU was provably halted

Observations Log

Per the handout (Section 3.7): things noticed while reading/running the code that look wrong, surprising, or undocumented, kept to the source we have (user-space + host server/). No fixes are made here — this milestone does not touch the code.

1. server/distinf_sim.py ring runs are flaky under WSL2 even with correct config

File/context: server/distinf_sim.py, server/mcast_switch.py (host-side multicast bridge), observed while bringing up the multi-worker ring.

Two distinct symptoms were seen, both during the shard/embedding fetch phase of a ring run, neither in the code itself:

- **Total failure** (never assigned layers within 900s, master log shows arp_lookup: resolution timed out for 0xa000001): the master could never reach the host-side weight server over the QEMU multicast segment at all, even though worker↔master registration over the *same* segment succeeded. Root cause: WSL2's default NAT networking mode does not forward IP multicast between host processes and QEMU guests. Fix: enable WSL2 **mirrored networking mode** (%USERPROFILE%\wslconfig, networking-Mode=mirrored, then wsl --shutdown).
- **Partial failure** (Transfer failed: received only N/72001 chunks after 6 rounds), seen even *after* the mirrored-mode fix, on both 2- and 3-worker rings: roughly 50-90 chunks (~0.1%) of the ~35 MB embedding transfer are lost and the transfer's 6-round retry budget is exhausted. This recurs intermittently — the identical command succeeds roughly half the time on a plain retry — so it looks like probabilistic UDP loss on the multicast segment under WSL2, not a deterministic defect. Higher worker counts (more co-resident QEMU VMs contending for the host's 8 cores) appear to make it more likely, but it is not exclusive to high worker counts.

```
baseline: X86_CPU=1 ./boot.sh, 3 prompts

xv6 kernel is booting

Scheduler: Round Robin (Default)
init: starting sh
$ llama -n 32 -i "Once upon a time" -t 0
No cached data found for llm_weights. Fetching from server...
PASS   llm_weights fetched and cached in shared memory (ID: 0).

No cached data found for llm_tokenizer. Fetching from server...
PASS   llm_tokenizer fetched and cached in shared memory (ID: 1).

thread_create: sharing user memory with main process (pid=3)
thread_create: sharing user memory with main process (pid=3)
thread_create: sharing user memory with main process (pid=3)
Once upon a time, there was a little girl named Lily. She loved to play outside in the sunshine. One day, she
saw a fox!

achieved tok/s: 3.594619

=====
FUNCTION PROFILING
=====
Function      Calls Total(ms) Self(ms) Avg(ms) Time(%)
-----
matmul         1 76618 76618.00 88.74%
fetch_model_weights 1376 8458 8493 6.24 9.83%
fetch_tokenizer 1 610 610 0.00 0.70%
forward        32 8837 255 276.15 0.29%
encode         1 173 173 0.00 0.20%
multithread_attention 192 58 62 0.30 0.07%
generate       1 9860 57 9860.00 0.04%
build_transformer 1 35 35 0.00 0.04%
rnnorm         416 19 38 0.04 0.03%
build_tokenizer 1 14 14 0.00 0.03%
sample         28 7 7 0.25 0.00%
decode         32 1 3 0.03 0.00%
build_sampler  1 2 2 2.00 0.00%
=====
Total Time: 86462 ms

=====
SYSTEM PERFORMANCE METRICS
=====

MEMORY USAGE:
Initial: 1.7 MB (1843200 bytes)
Peak: 66.5 MB (69763872 bytes)
Final: 66.5 MB (69763872 bytes)

TIMING:
Total (E2EL) 86462 ms
Inference 8624 ms

LUM METRICS:
TFTT 79369 ms

THROUGHPUT:
Generated 31 tokens
Rate 3.59 tokens/sec

=====
$ llama -n 32 -i "The quick brown fox" -t 0
PASS   Cached data found (segment: llm_weights, ID: 0). Ready for use.
PASS   Cached data found (segment: llm_tokenizer, ID: 1). Ready for use.
thread_create: sharing user memory with main process (pid=7)
thread_create: sharing user memory with main process (pid=7)
thread_create: sharing user memory with main process (pid=7)
The quick brown fox was walking through the forest. He was looking for something to eat. Suddenly, he saw a
big, juicy apple.

achieved tok/s: 2.633814

=====
FUNCTION PROFILING
=====
Function      Calls Total(ms) Self(ms) Avg(ms) Time(%)
-----
matmul         1376 11458 11515 8.32 94.33%
forward        32 11942 308 373.18 2.45%
encode         1 136 136 0.00 1.11%
multithread_attention 192 79 89 0.41 0.09%
generate       1 12142 51 12142.00 0.41%
rnnorm         416 25 42 0.06 0.34%
build_transformer 1 23 23 0.00 0.19%
fetch_model_weights 1 36 36 0.00 0.13%
fetch_tokenizer 1 12 12 0.00 0.09%
build_tokenizer 1 12 12 0.00 0.09%
sample         27 19 18 0.37 0.08%
decode         32 3 3 0.09 0.02%
build_sampler  1 1 1 1.00 0.00%
=====
Total Time: 12272 ms

=====
SYSTEM PERFORMANCE METRICS
=====

MEMORY USAGE:
Initial: 68.2 MB (63197184 bytes)
Peak: 66.5 MB (69795840 bytes)
Final: 66.5 MB (69795840 bytes)

TIMING:
Total (E2EL) 12272 ms
Inference 11778 ms

LUM METRICS:
TFTT 1819 ms

THROUGHPUT:
Generated 31 tokens
Rate 2.63 tokens/sec

=====
$ llama -n 32 -i "In the beginning" -t 0
PASS   Cached data found (segment: llm_weights, ID: 0). Ready for use.
PASS   Cached data found (segment: llm_tokenizer, ID: 1). Ready for use.
thread_create: sharing user memory with main process (pid=11)
thread_create: sharing user memory with main process (pid=11)
thread_create: sharing user memory with main process (pid=11)
In the beginning of the village, there was a little girl named Lucy. She was three years old and loved to
explore. One day, she was walking through

achieved tok/s: 4.266446

=====
FUNCTION PROFILING
=====
Function      Calls Total(ms) Self(ms) Avg(ms) Time(%)
-----
matmul         1376 7805 7824 5.14 92.55%
forward        32 7433 218 232.28 2.83%
encode         1 149 149 0.00 1.93%
multithread_attention 192 58 68 0.30 0.77%
generate       1 7632 34 7632.00 0.44%
rnnorm         416 20 31 0.04 0.40%
build_transformer 1 21 21 0.00 0.27%
fetch_model_weights 1 16 17 0.00 0.22%
fetch_tokenizer 1 15 15 0.00 0.19%
build_tokenizer 1 11 11 0.00 0.14%
sample         29 7 11 0.24 0.14%
decode         32 3 5 0.09 0.06%
build_sampler  1 1 1 1.00 0.01%
=====
Total Time: 7752 ms

=====
SYSTEM PERFORMANCE METRICS
=====

MEMORY USAGE:
Initial: 68.3 MB (63234048 bytes)
Peak: 66.6 MB (69849888 bytes)
Final: 66.6 MB (69849888 bytes)

TIMING:
Total (E2EL) 7752 ms
Inference 7266 ms

LUM METRICS:
TFTT 1373 ms

THROUGHPUT:
Generated 31 tokens
Rate 4.26 tokens/sec

=====
$ OPU: Terminated
```

Figure 19: Raw console output backing the single-node baseline table

```
baseline: distinf_sim.py --workers 3 --model 1 --steps 32

[sim] staging 4 prebuilt node image(s)...
[sim] starting host mcast switch (weight server on the segment)...
[sim] booting 4 node(s) (timeout 280s each)...
[sim] master: expecting 3 worker(s), model 1 (stories15M), 32 tokens
[sim] OK worker1 registered (1.4s)
[sim] OK worker2 registered (1.9s)
[sim] OK worker3 registered (2.7s)
[sim] OK 3 worker(s) registered (serialized)
[sim] OK layers assigned
[sim] OK worker1 shard resident (34s)
[sim] OK worker3 shard resident (34s)
[sim] OK worker2 shard resident (34s)
[sim] OK all shards resident -- generating

[sim] ==== generated ====
Once upon a time , there was a little girl named L ily . She loved to play outside in the sun sh ine . One day
, she saw a big
[sim] =====
[sim] ===== run summary =====
[sim] topology      master 10.0.0.2 + 3 worker(s) | model 1 (stories15M) | 32 steps
[sim] build         4 image(s) in 0s
[sim] boot          master 16s | worker1 16s | worker2 17s | worker3 17s
[sim] registration  worker1 1.4s | worker2 1.9s | worker3 2.7s   (serialized)
[sim] layers        w1 [0,2) | w2 [2,4) | w3 [4,6)
[sim] shard fetch   worker1 34s | worker3 34s | worker2 34s
[sim] final ring    3/3 worker(s)
[sim] generation    32 token(s) in 93.3s
[sim] latency       latency/token (avg over 32): ring rtt 2505960 us = compute 382276 us + network 2123684 us
| master 258655 us | 0.3 tok/s
[sim] rpc drops     malformed 0 | rate-limited 0 | bad-prog 0 | unknown-proc 0
[sim] logs          /tmp/distinf_sim/{master,worker1..N}.log
[sim] result        PASS
[sim] =====
[sim] PASS
```

Figure 20: Raw console output backing the 3-worker ring baseline row

Why this belongs in the log rather than “fixed and moved on”: this is a testing infrastructure/host-environment issue, not a defect in user/, server/, or the kernel under test — a marker running the identical runbook on a different host (native Linux, no multicast virtualisation layer in the way) would likely not see it at all. It cost real time to characterise, though, and directly explains why server/distinf_regression_test.py’s two ring-based tests (test_ring_generates, test_matches_single_node) were run three times on this machine: FAIL / FAIL / FAIL, all with the identical never assigned layers within 900s signature, while a manual distinf_sim.py --workers 2 CLI invocation outside pytest succeeded cleanly once (the transcript used for the Section 3.1.2 bring-up deliverable) and failed on other attempts. The result is inconsistent across identically-argued runs on the same machine within the same session — the clearest possible signal that the variable is host/network timing, not the code path, since the same code path produced a full correct token stream at least once. Exactly the “anything other than the expected counts is worth a log entry” case the handout calls out for this suite; the suite’s expected result table (2 passed, 6 skipped) could not be reproduced on this host despite three attempts, so the submitted evidence for distinf_regression_test.py is a failing run with this explanation attached, not a fabricated pass.

2. worker_entry_t.last_seen_ms (user/discovery.h:116) holds ticks, not ms

Files: user/discovery.h:116, user/distinf.c:1023,1027,1048.

The field is named last_seen_ms and every timeout constant it is compared against (SUSPECTED_TIMEOUT, EXPIRED_TIMEOUT in discovery.h:24-25) is built from HEARTBEAT_INTERVAL_TICKS, but the value actually stored in it is uptime() — xv6’s raw timer-tick counter, passed straight through as the now_ms parameter at every call site (discovery_handle_call(&msg, &src, uptime()), discovery_expire_workers(uptime()) in user/distinf.c). The unit is ticks throughout; the _ms in the name and the parameter name now_ms both claim a millisecond value that is never actually computed anywhere. Harmless only because every comparison uses the same (wrong) unit consistently — it would become a real bug the moment any code compared this field against an actual millisecond value (e.g. a wall-clock timestamp from a different source).

3. discovery_expire_workers ages a PENDING worker exactly like an ACTIVE one — the shipped diagrams/state_diagram.png doesn’t show this and is wrong

File: user/discovery.c:769-800, comment at 776-782 (the code documents this itself). See docs/figures/f3_worker_lifecycle.md for the full corrected state diagram and divergence paragraph (Section 3.4, figure F3).

4. handle_auth_hello’s “slot stays PENDING and ages out” comments (user/discovery.c:514,562) rely on behaviour that isn’t obviously intended

Both comments describe a PENDING worker that fails registration (no entropy for a probe nonce, or a bad CAP_ACK) as one that “ages out” — but per observation #3, aging out from PENDING was never a deliberately designed path; it is a side effect of discovery_expire_workers not special-casing PENDING. The comments read as if this were the intended behaviour, when the code’s own later comment (#3) calls it a known divergence from the design (“documented as a follow-up in README.md, not fixed here”). Not a

functional bug — the slot does get reclaimed — but the intent behind it is inconsistently described across the file.

5. `rdcycle/rdinstret` do not measure guest execution under this QEMU configuration — they track host wall-clock time almost exactly

Files: `kernel/syscall.h:39-41` (`SYS_rdcycle/SYS_rdtype/SYS_rdinstret`), `boot.sh/boot-gdb.sh` (the `qemu-system-riscv64` invocation — no `-icount` flag).

Verified directly with `gdb-multiarch` connected to `boot-gdb.sh`'s `gdb` stub (port 26000) *before ever sending continue* — i.e. with the guest vCPU provably halted (`-S`) and zero instructions executed. `info all-registers` exposes `cycle`, `instret`, `mcycle`, `minstret` as readable pseudo-registers. Sampled twice, 5 seconds of real time apart, with the vCPU never released from halt:

Sample	instret	cycle
t=0s	1,635,762,915,193	1,635,763,713,343
t=5s	1,647,061,538,553	1,647,062,338,672
Δ over 5s	11,298,623,360	11,298,625,329
implied rate	≈ 2.2597 GHz	≈ 2.2597 GHz

Both counters advanced by essentially the same amount, at a rate matching a modern host CPU's clock, **while the guest hart executed zero instructions**. This is only possible if QEMU is not gating these CSRs on actual guest instruction retirement — under TCG emulation without `-icount` (not passed by either boot script), `cycle/instret` appear to be implemented as `host_wallclock × constant`, not true per-instruction hardware counters. The two counters also track each other almost exactly (implied IPC ≈ 1.00 throughout, including while genuinely idle), which is itself further evidence they are the same underlying wall-clock-derived value read twice rather than independent hardware events.

Consequence for Section 3.6: the assignment's own guidance ("QEMU does not model real hardware timing... report ratios and relative differences") already anticipates some imprecision, but this goes further — `rdcycle/rdinstret`, as actually implemented in this environment, are not a higher-fidelity alternative to the 1 ms-resolution wall-clock timer at all; they are the *same* wall-clock signal at higher numeric resolution. The baseline measurements report derives `cycle/instret` figures from each run's precisely measured wall-clock duration and this empirically established ~ 2.2597 GHz scale factor (method documented in the baseline measurements section) rather than claiming a live per-run hardware sample, since a live sample would carry no more information than the wall-clock time already gives.

(Further entries to be added as the remaining subsystems — RPC retry logic, XDR bounds checks, shard layer-range arithmetic — are read in detail for the architecture map and figures.)

Related Systems (Section 3.8)

For each system: what it does that this project does not, and what this project does that it does not.

Stock xv6-riscv

What stock xv6 does that this system does not: ships as buildable source for its entire kernel (this project's kernel .c files are withheld for this milestone by design, not by the upstream project); stays minimal and single-purpose as a teaching OS with no networking, no threads, no shared memory, and no resource-limit/capability model — it is deliberately *not* trying to run a real workload.

What this system does that stock xv6 does not: a full network stack (ARP/IP/UDP, kernel/net.c), a from-scratch RPC/XDR layer (RFC 5531/4506), a distributed application (pipeline-parallel LLM inference) running across multiple instances of itself, threads and shared memory (thread_create, shmget/shmat), and a real hardening set — stack canaries with genuine boot entropy, ASLR, POSIX resource limits, a capability bitmask, and W^X enforcement — none of which stock xv6 has.

Linux

What Linux does that this system does not: true multi-user isolation (full uid *and* gid semantics, file ownership/mode bits, chown/chmod), full POSIX.1e capability *sets* (this system has one effective bitmask covering a single capability, not the ~40 distinct Linux capabilities), syscall-filtering sandboxes (seccomp/pledge-equivalents), cgroups-based resource accounting across many dimensions (this system has exactly two limited resources, RLIMIT_AS/RLIMIT_CPU), loadable kernel modules, SMP scheduling at real scale, and a production TCP stack (this system's transport is UDP-only — RPC record marking over TCP, RFC 5531 §11, is explicitly unimplemented, per user/rpc.c's header comment).

What this system does that Linux does not: nothing Linux *can't* do in principle — the comparison that matters here is legibility, not capability. Every layer of this system, kernel included, is small enough to read end-to-end and reason about exhaustively (the whole point of building the distributed-inference workload on xv6 rather than on Linux, per the project's stated rationale); a security claim here (the W^X check, the capability-probe ceiling, the HMAC identity check) can be pinned down to a specific function and line, which is infeasible to do with the same confidence against the Linux kernel and glibc.

xv6-llm-runtime-arch (github.com/syedtaha22/xv6-llm-runtime-arch — the base this project builds on)

What it does that this system does not: establishes the baseline runtime architecture and single-node inference path this project extends — the starting point, not a competing design.

What this system does that it does not: everything in the L2/L3 layers this project adds on top of that base — the hardened network stack (ARP/IP/ICMP/UDP filtering, rate limiting, bogon/ingress checks), the RPC/discovery/ring layer that turns a single-node runtime into a pipeline-parallel *distributed* one, the PSK/HMAC worker authentication, the sized capability probe, and the kernel hardening set (canaries, entropy, resource limits, capabilities, W^X) — none of which the base runtime architecture includes; it is a single-node inference engine, not a distributed, adversarially-tested one.

Exo (github.com/exo-explore/exo) — closest prior-art pipeline-parallel inference system

What Exo does that this system does not: runs on real heterogeneous consumer hardware (phones, laptops, Raspberry Pis) at real model scale, with a far more capable scheduler that adapts to heterogeneous device throughput and memory, and a polished user-facing API — none of which this project attempts; this project targets a from-scratch teaching kernel at toy model scale (15M/110M-parameter checkpoints) as a research testbed, not a usable inference product.

What this system does that Exo does not: authentication on peer discovery (Exo’s peer join has none), encryption/integrity on the wire (this system’s RPC is plaintext but HMAC-authenticated at registration and its packets get real input validation), and input validation on incoming tensors at the network/RPC boundary. This is exactly the gap this project’s own documentation identifies as its reason for choosing this workload: Exo is undefended against the same threat model (capability lying, identity spoofing, result corruption) that this project’s L3 layer explicitly targets and partially closes — including the caveat, established by this project’s own adversarial testing, that result corruption (a node that computes honestly but reports a tampered value) is *not* currently detected here either, so the comparison is “closes more of the gap than Exo does,” not “fully solved.”

Other pipeline-parallel inference systems (general class: Petals, vLLM’s pipeline-parallel mode)

What they do that this system does not: operate at production model scale (tens to hundreds of billions of parameters) with batching, continuous scheduling across many concurrent requests, and tensor parallelism in addition to pipeline parallelism (this project’s own README states tensor parallelism is out of scope for v1) — throughput and utilization at that scale is simply not this project’s goal.

What this system does that they do not: runs the *entire* stack, network driver through inference kernel, inside a from-scratch, auditable teaching OS with no underlying Linux/CUDA/production-scheduler dependency — every hop’s authentication, every packet’s validation, and every allocation’s resource accounting is enforced by code this project’s own team can read in full, which is not true of any production system built on Linux and a mainstream ML runtime.

Contributions (Section 3.9, bonus)

What we believe is worth merging, and the evidence:

1. **F3's divergence identification.** `diagrams/state_diagram.png` is wrong in a way the code itself flags (`discovery.c:777-782`) but that has apparently never been turned into a corrected diagram. `docs/figures/f3_worker_lifecycle.md` supplies one, derived line-by-line from `discovery.c/discovery.h`, and is a drop-in replacement for the shipped PNG.
2. **The `rdcycle/rdinstret` wall-clock-equivalence finding** (Observations Log #5). This isn't just a note — it changes how baseline measurements (and any future milestone's performance work) should be reported: a live per-run `rdcycle/rdinstret` sample under this QEMU configuration carries no information beyond the wall-clock timer already available, so future measurement work can skip instrumenting the counters directly and use wall-clock time (scaled by the empirically-derived ~ 2.2597 GHz constant if a cycle-shaped number is wanted for presentation) without loss of fidelity. This is evidenced with a reproducible, from-scratch experiment (`docs/transcripts/cycle_instret_probe.txt`) any future milestone team can rerun in under 10 seconds to confirm.
3. **The WSL2 mirrored-networking requirement**, now documented as an explicit runbook prerequisite (`docs/figures/runbook.md`). Undocumented, this cost real time (multiple 900-second timeouts before the root cause was found) and would cost the same to any other team member or grader running the ring on WSL2 for the first time.
4. **Evidence that `distinf_regression_test.py`'s ring-dependent tests are environment-flaky, not code-broken, on this host class.** Three consecutive identical invocations produced FAIL/FAIL/FAIL with the same signature, while a manual `distinf_sim.py` invocation with the same arguments succeeded once cleanly and reproduced the single-node token stream exactly. This distinguishes “the pipeline is broken” from “this specific CI-style gate is unreliable on a resource-constrained virtualised host,” which matters for how much weight Milestone 2/3 should put on this particular gate's pass/fail signal on similar hardware.

Appendices

AI Use Disclosure

This milestone's environment setup, system bring-up, reverse-engineering (warm-up call graph and kernel gdb trace), architecture mapping, all seven Mermaid figures, the syscall diff, the baseline-measurement methodology and its execution, the observations log, the related-systems comparison, and the drafting of this report were carried out using **Claude Code** (Anthropic), operating the toolchain (QEMU, gdb, pytest, LaTeX, mermaid-cli) directly inside a WSL2 Ubuntu environment under the direction and supervision of the team. Every quantitative claim in this report (regression suite pass counts, baseline timings, syscall numbers, line/address citations) was produced by actually running the referenced command against this repository's shipped binaries and source, not generated from general knowledge - the raw transcripts backing each claim are committed alongside this report under docs/transcripts/ and docs/screenshots/.

Repository Layout of Supporting Evidence

- docs/figures/ - Markdown source for every report section, and mmd/ holding each Mermaid diagram's standalone source file.
- docs/screenshots/ - rendered PNGs: all seven figures plus every console transcript, rendered as a terminal-styled image via docs/tools/termshot.py.
- docs/transcripts/ - raw, unedited console/pytest/gdb output for every run referenced in this report.
- docs/tools/ - termshot.py (transcript-to-image renderer) and gdb_trace_cmds.gdb (the batch script used for the kernel reverse-engineering trace).

Commit Log

See the repository's git log on <https://github.com/Arjcodes123/OS-Milestone-1> for the commit history backing this submission.